

```
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, roc_auc_score, confusion_matrix
from sklearn.neighbors import KNeighborsClassifier
from sklearn.over_sampling import SMOTE
import seaborn as sns
import matplotlib.pyplot as plt
import xgboost as xgb

# Combine datasets
df = pd.read_csv('data/ALLFLOWMETER_HIKARI2021.csv')

# Display the first few rows
df.head()

# ...

# Compute the correlation matrix
correlation_matrix = data_cleaned.corr()

# Plot the heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=False, cmap='coolwarm')
plt.title('Feature Correlation Matrix')
plt.show()

# ...

# Define models
log_reg = LogisticRegression(max_iter=1000)

rf = RandomForestClassifier()

gbm = GradientBoostingClassifier()

xgb_model = xgb.XGBClassifier()

log_reg.fit(X_train, y_train)
y_pred_log_reg = log_reg.predict(X_test)

rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)

gbm.fit(X_train, y_train)
y_pred_gbm = gbm.predict(X_test)

xgb_model.fit(X_train, y_train)
y_pred_xgb = xgb_model.predict(X_test)

# Evaluation Metrics Function
def print_evaluation_metrics(model_name, y_test, y_pred):
    print(f'Evaluation Metrics for {model_name}:')
    print(f'Accuracy: {accuracy_score(y_test, y_pred):.4f}')
    print(f'Precision: {precision_score(y_test, y_pred):.4f}')
    print(f'Recall: {recall_score(y_test, y_pred):.4f}')
    print(f'F1-Score: {f1_score(y_test, y_pred):.4f}')
    print(f'ROC-AUC: {roc_auc_score(y_test, y_pred):.4f}')
    print("\n")

print_evaluation_metrics("Logistic Regression", y_test, y_pred_log_reg)

print_evaluation_metrics("Random Forest", y_test, y_pred_rf)

print_evaluation_metrics("Gradient Boosting Machine (GBM)", y_test, y_pred_gbm)

print_evaluation_metrics("XGBoost", y_test, y_pred_xgb)

# Perform cross-validation
cv_scores_log_reg = cross_val_score(log_reg, X, y, cv=5, scoring='accuracy')
cv_scores_rf = cross_val_score(rf, X, y, cv=5, scoring='accuracy')
cv_scores_gbm = cross_val_score(gbm, X, y, cv=5, scoring='accuracy')
cv_scores_xgb = cross_val_score(xgb_model, X, y, cv=5, scoring='accuracy')

# Print cross-validation scores
print("Cross-Validation Scores:")
print(f"Logistic Regression: (cv_scores_log_reg.mean():.4f) ± (cv_scores_log_reg.std():.4f)")
print(f"Random Forest: (cv_scores_rf.mean():.4f) ± (cv_scores_rf.std():.4f)")
print(f"Gradient Boosting Machine (GBM): (cv_scores_gbm.mean():.4f) ± (cv_scores_gbm.std():.4f)")
print(f"XGBoost: (cv_scores_xgb.mean():.4f) ± (cv_scores_xgb.std():.4f)")

/Users/sako/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(
/Users/sako/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(
/Users/sako/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(
/Users/sako/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(
/Users/sako/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(
/Users/sako/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_1 = _check_optimize_result(
/Users/sako/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Logistic Regression: 0.9017 ± 0.0180
Random Forest: 0.8659 ± 0.0740
Gradient Boosting Machine (GBM): 0.8900 ± 0.0341
XGBoost: 0.8761 ± 0.0574

# Feature Importance for Random Forest
rf.feature_importances_ = pd.DataFrame(rf.feature_importances_, index=X.columns, columns=['Importance']).sort_values(by='Importance', ascending=False)
plt.show()

gbm.feature_importances_ = pd.DataFrame(gbm.feature_importances_, index=X.columns, columns=['Importance']).sort_values(by='Importance', ascending=False)
gbm.feature_importances_.head(10).plot(kind='bar', title='Top 10 Feature Importances from Random Forest')
plt.show()

# ...

# Top 10 Feature Importances from Random Forest
plt.show()

# Top 10 Feature Importances from GBM
plt.show()
```